

Supervised Deep Learning

Marco Ballarin 1228022
marco.ballarin.6@studenti.unipd.it

ABSTRACT

In this report we will review the regression task on a polynomial function of unknown order and a classification task on the MNIST dataset. We will review with particular attention the use of optimizers, regularization methods, hyperparameters optimization and feature visualization for the networks.

1 Introduction

In this report we will tackle two different problems using the supervised deep learning framework. This means that we will have datasets D formed of couples $\{(x_i, y_i)\}_{i=1, \dots, N}$, where x_i are the samples and y_i the labels. The report will be divided into four main Sections:

1. Introduction, where we introduce our work;
2. Regression, where through a deep neural network we try to predict the behavior of a curve;
3. Classification, where through a convolutional neural network we try to classify the handwritten digits from 0 to 9 using the MNSIT dataset;
4. Appendix, where we will add images that are not fundamental for the flow of the report, but may be useful.

2 Regression

2.1 Methods

For the regression task it is better to use a reduced number of hidden layers, and in particular we choose to use 2 of them. Our network will so be composed of:

- An input layer of one neuron;
- A first hidden layer oh N_h neurons;
- A second hidden layer of $2 \cdot N_h$ neurons;
- An output layer of one neuron.

To avoid vanishing problems in the gradient we will use a rectified linear unit (ReLU) as activation function. This simple topology in the network is due to the fact that this problem is "simple" from the neural networks point of view, and adding too many parameters we will be easily led to overfitting. We will so make an extensive use of regularization procedures to avoid these problems. In particular, we will use two main methods:

1. dropout, which consists of randomly eliminate neurons during the training with probability p_{drop} .
2. L2 regularization, which consists of adding a penalty in the loss proportional to the 2-norm of the parameters, i.e. weights and biases. We stress that in pytorch this method is enabled not in the loss function but in the optimizer, using the `weight_decay` keyword.

However, implementing both forms of regularizations generated sub-optimal results. We so decided to put to 0 the L2 regularization, stressing that we can do this without losing generality.

We will try two different optimizers:

1. The stochastic gradient descent with momentum, which is the really basic algorithm for optimization. The addition of the momentum, i.e. the memory of previous steps, really helps to go out of local minima and speeds up convergence;
2. the Adam algorithm, adaptive moment estimator, a more advanced algorithm that uses estimations of first and second moments of the gradient to adapt the learning rate of each weight of the neural network.

The amount of data available for this task is quite limited, we will so apply a cross-validation technique to perform the hyper parameters tuning. This method consists of splitting the training test into k different folds and then, for each hyper parameters set, train the network on the union of $k - 1$ folds, performing the validation on the one which is missing. This approach ensure also better performances, since the validation set is different for each of the iterations, and so we cannot end up with a model that perfectly fits one given validation set. After the hyper parameters choice is done, we will train the model over the full training test.

We will then optimize the number of hidden units in the layers of the network, the learning rate and the number of epochs needed for the training procedure.

Lastly, we will visualize the network features: in particular we will plot the weights histograms and the activation profiles. We will see that we can actually extract a lot of information, in particular from the activation profiles.

2.2 Results

The best performing network is reported in Figure 1. The grid search lead to the parameters in




Figure 1: Final network for the regression task. The activation function is the rectified linear unit.

Table 1.

Dropout	Epochs	Learning rate	Optimizer	Regularization	Hidden units
0.023	400	0.0004	Adam	0	75

Table 1: Optimal parameters for the regression task.

We can finally test our network on the test set. We stress that the splitting of test and training test was not optimal, since we have two missing intervals in the input space. However, it is not always possible to have the full dataset, and we so pushed for the generalization capabilities of the network, using in particular the regularization which mitigates the overfitting problem. We can see from Figure 2 that the first maximum is well predicted by the model, while the second one is cut. To gain a better insight in why we have this behavior we investigate the weight histograms and the




Figure 2: Data points with the model prediction. We can see that the model is able to reproduce the first maximum, but cuts the second.

activation profiles of the network. From Figure 10 we can not really extract useful information. It is instead really important to observe the activation profiles, in which we plot on a 3d plot the activation profile wrt the input space for some of the neurons. In the first layer, presented in Figure 3, we can see that the network strongly divide the right and left part of the input space, as we can expect since the behavior in the two regions is completely different. The second layer profile, in Figure 4, shows instead a profile more similar to the output of the network. We can see that some neurons actually have located the second maximum of the function, but on average it is hidden, as we can see from the thick red line.

We can so conclude that we had an acceptable result in the regression task, even if not perfect.

3 Classification

3.1 Methods

In the classification task we have to develop a network which classifies the MNIST dataset, i.e. the dataset containing handwritten digits from 0 to 9. This dataset is quite exhaustive, and so we will not need to apply the Cross-validation technique. The size of this dataset present nevertheless another challenge: the training time can become too long. For this reason we really need to finely optimize the number of epochs we use in this procedure. One possible solution would be, as in the previous case, use them as a hyper parameter. However, we will waste a lot of computational resources in this search. We so decided to apply a technique that is called early stopping: if the validation loss l_v does not diminish sufficiently in a given number of epochs then we kill the training loop. In this way we can get rid of poorly initialized networks and have an overall speed up of the process. In particular, calling $\vec{l}^{(t)} = (l_v^{(1)}, l_v^{(2)}, \dots, l_v^{(t)})$ the vector containing all the losses at iteration t , with a maximum number of iteration T we stop the process if:

$$\bar{l} - l_v^{(t)} < 0.01 \quad \text{with} \quad \begin{cases} \bar{l} = \frac{1}{t-\tau} \sum_{i=\tau}^t l_v^{(i)} \\ \tau = \min(t, T/10) \end{cases} \quad (1)$$

The other main difference with respect to the regression task is that the dataset is composed by images. For this reason we decided to use two convolutional layers at the beginning of the network.




Figure 3: Activation profile of the first hidden layer. In thick red we can see the average of the profiles of the neurons, multiplied by a constant to better catch the behavior. It is clear that this layer is able to understand if the input is on the left or right side of the input domain.




Figure 4: Activation profile of the second hidden layer. The average enhanced activation is defined as in Figure 3. We can notice from this plot that the first maximum is well represented, while the second one is cut even if some of the neurons actually are particularly active in that interval.

These layers make use of filters, i.e. feature maps that are applied along all the image with the same weights. We will use filters of dimension 3×3 and 5×5 . This means that we will employ several small $n \times n$ feature maps that will enhance the same details along all the figure. The remarkable example in this framework are filters which enhance edges. We will analyze them in Section 3.2.

Through a random search we will optimize the following parameters:

- The dropout;
- The learning rate;
- The optimizer, choosing from stochastic gradient descent and Adam;
- the strength of the L2 regularization,

To further enhance the dataset we will apply a random rotation of 45° and some Gaussian noise. To help the network in the learning phase we will also normalize the inputs, using the apposite pytorch transformation. We present the network structure in Figure 5

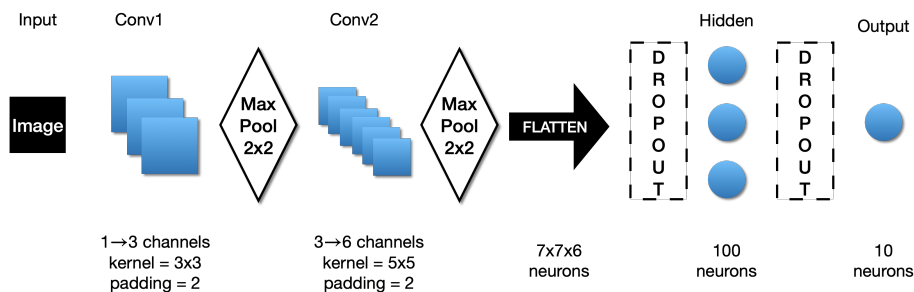


Figure 5: Network structure for the classification task. The activation function used is the rectified linear unit. The output digit is selected as the argmax among the 10 output neurons.

3.2 Results

We present the grid search space in Figure 11, and report the best parameters found in Table 2. With this set of parameter the network results fully trained in only 7 epochs. We can see that the


Dropout	Learning rate	Optimizer	Regularization
0.1	0.003	Adam	0.001

Table 2: Optimal parameters for the classification task.

regularization, through both the dropout and the L2 regularization, is higher in this case, meaning that this task requires particular attention to avoid overfitting.

The accuracy on the test set of the model gave optimal results, with an accuracy of 0.999.

We present in Figure 6 the filters of the first convolutional layer and in Figure 7 the filters of the second convolutional layer. They are, however, of difficult interpretation. We can however see in Figure 7 a pattern that suggests that these filters are looking for edges in the figures.



It is more interesting to analyze the activation profiles of the convolutional layers. In Figure 8 we can observe in a really clear way the digit that was given in input, while in 9 we can observe the enhancement of the edges of that number, as suggested by the filter.

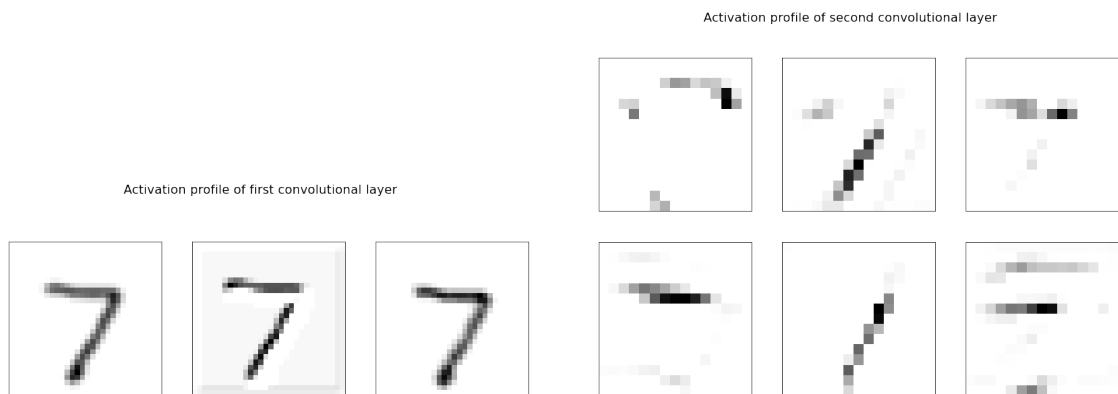


Figure 8: The three channels in output from the first convolutional layer. We can observe the input digit, which is a 7.

Figure 9: The six channels in output from the second convolutional layer. We can observe a decomposition of the input digit in edges, which is common in convolutional networks.

4 Appendix

We present here images that can be interesting but are not fundamental for the report.




Figure 10: Weight histograms for the different layers of the network.




Figure 11: Hyperparameters search for the classification task. We plotted with a color which is proportional to the loss function, reported in log-scale. The darker the color, the better the result. We notice that a really high L2 regularization usually translates in a low loss score.